

Scalable Retail Inventory Analytics Using Big Data Technologies

Evan Banks

College of Computing and Software Engineering

Kennesaw State University

Marietta, Georgia, USA

ebanks28@students.kennesaw.edu

Abstract—This project proposes a scalable retail inventory analytics system using distributed storage and parallel data processing. The system will be designed to take in large-scale inventory datasets, store them in a distributed file system, and support analytical queries using distributed processing frameworks. This project covers Big Data architecture design, scalability, and multiple layers of the Big Data technology stack.

I. PROJECT OVERVIEW

A. Project Title

Scalable Retail Inventory Analytics Using Big Data Technologies

B. Domain

E-Commerce, specifically, retail and supply chain analytics.

C. Problem Statement

Retail businesses generate large volumes of inventory and sales data across products, stores, and time periods. Analysis of this data is important to understand inventory levels, identify demand patterns, and reduce issues such as overstocking and understocking. The problem comes from the discrepancy between the growth rates of data volume and storage capacity of a single machine. As data volume increases, single-machine data processing becomes inefficient.

This project confronts the challenge of performing scalable inventory analytics by designing a distributed Big Data system capable of reading, storing, processing, and querying large-scale retail inventory datasets. The goal of this system is to support analytical queries such as inventory aggregation, trend analysis across seasons, and historical comparisons using parallel processing and distributed storage.

D. Scope

The scope of this project is limited to batch-oriented analytical processing of retail inventory data using Big Data technologies. The system will focus on exploratory data analysis and combined queries rather than real-time processing or production-level deployment.

This project does not include real-time streaming ingestion, advanced machine learning models, or deployment in large commercial cloud clusters. Instead, it emphasizes system architecture design, engagement with multiple layers of the Big Data stack, and the application of distributed storage and parallel processing concepts.

II. SYSTEM DESCRIPTION

A. Data Sources

The main data source used in this project is retail inventory and sales data. This data contains product stock and sales data, as well as time-related attributes such as dates and seasons. The data can be obtained from publicly available retail data sources or generated synthetically to represent realistic retail data.

The data will be stored in flat files and processed in batch mode in the distributed storage layer. The data can be generated synthetically to make it larger and represent multiple retail stores, products, and seasons to necessitate the need for distributed processing.

B. Data Characteristics

Volume: The size of the dataset is expected to vary from a few gigabytes to tens of gigabytes, which will encompass historical sales and inventory data for various products and stores. The size of the dataset is expected to be larger than the capacity of a single machine.

Variety: various structured and semi-structured fields such as product codes, store information, inventory, sales data, timestamps, and season variables. The input data will be stored in CSV files and may be converted to columnar storage formats like Parquet for efficient analytical processing.

Velocity: The data ingestion process is based on a batch processing paradigm, where the inventory and sales data are loaded periodically. The project does not target real-time processing but rather focuses on batch analytics for large datasets.

C. Stack Layers

This project utilizes multiple layers of the Big Data technology stack, namely storage, processing, and querying. Each layer is selected to address scalability, fault tolerance, and efficient logical processing in large inventory datasets.

Storage Layer: The storage layer will be implemented using the Hadoop Distributed File System (HDFS). HDFS stores large files by dividing them into fixed-size blocks distributed across multiple nodes. This enables parallel data access and fault tolerance through replication. This design allows the system to scale storage capacity horizontally while keeping data available even when nodes fail.

Processing Layer: Apache Spark will be the main data processing framework. Spark supports parallel data transformations through Resilient Distributed Datasets (RDDs) and higher-level DataFrame abstractions. Its distributed execution model enables efficient aggregation, filtering, and joining of large datasets, while in-memory processing boosts performance for iterative analytical workloads.

Query Layer: The query layer will be established using Apache Hive or Spark SQL, both of which provide declarative SQL-based interfaces for querying data stored in HDFS. Queries are executed in a distributed manner across the cluster, allowing analytical operations such as grouping, aggregation, and trend analysis to scale with data volume. Users can analyze inventory data without needing low-level programming constructs.

D. Assumptions

The design of this project is based on several assumptions. The system will be developed and evaluated in a controlled environment using a pseudo-distributed or small-scale cluster configuration rather than a full production cluster. Data is assumed to be static during processing, with no concurrent updates or real-time data streams.

Inventory and sales datasets are assumed to be well-structured and free of significant data quality issues beyond basic pre-processing requirements. In addition, system security, access control, and fault recovery beyond the default framework mechanisms are outside the scope of this project.

III. IMPLEMENTATION APPROACH

A. Technology Choices

This project utilizes a combination of open-source Big Data technologies selected to address scalability, fault tolerance, and analytical processing requirements. Hadoop Distributed File System (HDFS) is used for distributed storage due to its ability to replicate data and provide high-throughput access to files.

Apache Spark is used as the main processing technology based on its ability to process data in parallel using in-memory computation. Apache Spark also has integration capabilities with SQL-based data analysis using Spark SQL.

Apache Hive is used as a metadata management tool that allows users to write queries in SQL, providing a standard interface for querying large datasets in HDFS.

B. Processing Model

This system uses a batch processing approach in which inventory and sales data is read, processed, and analyzed in a series of execution cycles. The data will be stored in the distributed storage layer as static data files, which will be processed by Apache Spark jobs to transform the data by performing operations such as filtering, aggregation, and joining.

The batch processing model is most appropriate for historical inventory analysis, where queries are focused on trends and summary statistics rather than real-time responses.

C. Scalability Plan

This system is designed to be horizontally scalable with the growing volume of data and the complexity of the analysis to be performed. Storage scalability is achieved through HDFS, which enables data to be distributed across multiple nodes and replicated with the growing size of the cluster to maintain fault tolerance.

Processing scalability is achieved with the Apache Spark parallel processing model, which enables concurrent execution of data processing tasks on multiple worker nodes by breaking down the data processing task into stages. With the growing volume of data, additional processing nodes can be added to the cluster for maintaining reasonable execution times for tasks.

Scalability is also achieved through data partitioning techniques based on product identifiers or store locations. The partitioned data enables concurrent execution and efficient query processing with reduced data scans.

D. Evaluation Metrics

The performance of the system will be measured using performance metrics pertinent to batch-based Big Data Analytics. Examples of performance metrics include job execution time, query response time, and the amount of data processed during analytical processes.

The performance evaluation may also be extended to include system scalability. This can be achieved by observing the execution time as the amount of data used for execution increases. These performance metrics will help in understanding how well distributed storage and parallel processing help in inventory analytics.

E. Project Repository

The source code and project materials for this work are maintained in a public GitHub repository available at: <https://github.com/ebanks28/cs4265-retail-inventory-big-data-project.git>

IV. LITERATURE REVIEW

A. Distributed File Systems: Google File System

Ghemawat et al. [1] describe the Google File System (GFS), a distributed file system that enables the storage and processing of large datasets on commodity hardware. The GFS is a solution to scalability, fault tolerance, and high-throughput data access issues by splitting files into fixed-size blocks and replicating them across multiple machines.

One of the major contributions of GFS is its focus on fault tolerance with replication and relaxed consistency requirements. By assuming frequent hardware failures, GFS enables the system to be available while allowing parallel access to data. This approach is the basis of the Hadoop Distributed File System (HDFS), which also uses a block-level storage solution with replication.

In this project, the ideas presented in GFS are implemented using HDFS as the main storage solution for retail inventory and sales data. Distributed block-level storage enables parallel processing using Apache Spark and allows the system to scale storage capacity with increasing data size. Replication

provides data availability for batch analytical processing, making distributed file systems an appropriate solution for large-scale inventory analytics.

B. Parallel Data Processing: MapReduce

Dean and Ghemawat [2] present the MapReduce programming model, which is a framework that facilitates data processing by using parallel processing techniques for data computation. MapReduce is a model that allows for parallel processing, where data is processed in parallel using *map* and *reduce* functions.

MapReduce facilitates scalable processing by dividing input data, processing it in parallel, shuffling data, and combining output using reduce functions. MapReduce is best suited for batch processing, where data is filtered, aggregated, and summarized.

Although MapReduce jobs are not implemented in this project, the underlying processing concepts are reflected in the system. Apache Spark is a generalized version of MapReduce, where execution is not limited to MapReduce. Although not explicitly implemented in this project, batch analytics performed on the retail inventory data, such as grouping inventory levels by product or store, are based on parallel processing concepts introduced by MapReduce.

C. In-Memory Distributed Processing: Apache Spark

Zaharia et al. [3] describe a distributed computing framework called Apache Spark, which aims to accelerate data processing on a large scale through the implementation of in-memory computing. The framework utilizes a data structure called Resilient Distributed Datasets (RDDs), a collection of data that can be processed in parallel across a cluster of nodes.

The major contribution of the Spark framework lies in its implementation of a data structure called a directed acyclic graph (DAG), which facilitates more flexible computing than the conventional MapReduce paradigm. The implementation of in-memory computing helps minimize disk I/O, thus increasing the overall performance of the data processing engine.

In this project, the Apache Spark engine acts as a major data processing engine to process huge amounts of data related to retail inventory, which is stored on HDFS. The parallel transformation capabilities of Spark are used to process the data, such as filtering, aggregating, and joining data related to inventory and sales.

D. Distributed Querying: Apache Hive

Thusoo et al. [4] describe Apache Hive, a data warehousing system developed using a distributed storage and processing infrastructure to support efficient analytical queries using a SQL-like language. Using Hive, users can define metadata over data stored in a distributed file system and execute declarative queries without having to write distributed code.

The major contribution of Hive lies in its support for a "schema-on-read" paradigm, which enables users to execute structured queries over large amounts of data stored in raw formats, such as CSV files, or even columnar storage

formats, such as Parquet files. Using HiveQL, users can define queries over data stored in a distributed environment that are then executed as a distributed plan to support efficient data aggregation, filtering, and joining over large data sets.

Hive and Spark SQL form the main querying interface for this project to support the querying of retail inventory and sales data stored in the HDFS environment. By providing a SQL-based querying interface, the system enables efficient data analysis tasks, including inventory aggregation and trend analysis, on data stored in a distributed environment.

E. Big Data Architecture and Stack Concepts

Fourny [5] provides a thorough discussion of Big Data systems by tracing the history of data storage, processing, and querying technology. The reading material introduces key ideas like the Three Vs of Big Data, distributed file systems, data models, and parallel processing frameworks, with a focus on the role of layered architectures in scalable data-intensive systems.

One of the main ideas explored in the textbook is the Big Data technology stack, which breaks down concerns into separate layers for storage, processing, and querying. Through the discussion of systems like HDFS, MapReduce, Spark, and Hive in a single architectural context, the reading material illustrates how distributed systems address scale, throughput, and fault tolerance concerns.

The system architecture developed in this project is a direct application of the ideas discussed in the Big Data Textbook. The layered architecture of distributed storage in HDFS, parallel batch processing in Apache Spark, and declarative querying in Hive or Spark SQL is consistent with the stack-based approach described in Fourny. This consistency ensures that the project architecture adheres to established Big Data best practices while meeting the analytical needs of retail inventory data at scale.

V. ARCHITECTURE DIAGRAMS

Figure 1 illustrates the layered Big Data architecture used in this project, highlighting the separation between storage, processing, and query components.

Figure 2 shows the batch-oriented data flow and processing pipeline used for inventory analytics.

REFERENCES

- [1] S. Ghemawat, H. Gobioff, and S.-T. Leung, "The Google File System," Proc. 19th ACM Symp. Operating Syst. Principles, 2003.
- [2] J. Dean and S. Ghemawat, "MapReduce: Simplified Data Processing on Large Clusters," Proc. 6th Symp. Operating Syst. Design and Implementation, 2004.
- [3] M. Zaharia et al., "Spark: Cluster Computing with Working Sets," Proc. 2nd USENIX Conf. Hot Topics in Cloud Computing, 2010.
- [4] A. Thusoo et al., "Hive: A Warehousing Solution Over a MapReduce Framework," Proc. VLDB Endowment, vol. 2, no. 2, 2009.
- [5] G. Fourny, *The Big Data Textbook: From Clay Tablets to Data Lakehouses*, 2nd ed., 2024.




Fig. 1. Stack architecture of the distributed retail inventory analytics system




Fig. 2. Data flow and processing pipeline for batch inventory analytics